

Vibhu Dixit

vibhu.dixit02@gmail.com | 623-299-0920 | tinyurl.com/vibhu-portfolio | [LinkedIn:vibhu-dixit-swe](#) | [GitHub](#)

Summary

Full-Stack Software Engineer with focus on high-velocity product delivery. Experienced in building end-to-end AI applications. Proven track record of operating autonomously and deploying production-ready integrations. Looking to join engineering focused company where I can work on impactful products, and continue growing as a software engineer.

Skills

Languages: JavaScript, TypeScript, Python, SQL

Frontend Frameworks: React, Next.js

Backend & APIs: Django, FastAPI, NestJS, REST

Databases & Data: PostgreSQL, MySQL, Oracle, Neo4j, Redis, Apache Kafka, Airflow, Iceberg, ETL/ELT

Cloud Infrastructure & AI: AWS, Azure, Docker, Linux, CI/CD, GitHub Actions, Jenkins, LangChain, LangGraph

Professional Experience

Software Engineer (AI) Intern — Privv Inc.

Jan 2026 – May 2026

- Engineered a bulk invoicing module using **AWS Textract** as OCR tool, enabling document extraction and automation that reduced manual accounting overhead by 80% across vendor and contract workflows for project managers.
- Architected the core engine for agentic customized reporting, defining system design and data flow for a flagship feature in a capital project management platform, leading to faster decision making.
- Built modular **RESTful microservices** and accompanying unit/integration test suites, improving reliability and production stability for AI-driven features used in real-time contractor decision-making.

Data Engineer Intern — Alcon Inc.

Feb 2024 – Aug 2024

- Optimized data persistence layers in DynamoDB for the **Inventory Analytics Dashboard**, testing different keys for partitioning, improving query performance and data visibility for by **40%**.
- Designed high-availability **Airflow ETL pipelines** for the Global Telemetry Platform, processing 50M+ events daily to enable real-time health monitoring of surgical equipment.
- Containerized Airflow triggered ETL worker services using Docker, ensuring consistent deployment of the Data Ingestion Service across development and production environments.

Projects

Hotel Booking Agent [\[link\]](#)

March 2026 - April 2026

- Built Linq webhooks and SMS trip extraction flows using single JSON calls per turn (12k-char input, 2048-token output), converting unstructured text into structured trip state with lower latency.
- Engineered a voice pipeline using STT and relevance filtering to extract travel entities before mapping, reducing token payloads and improving structured extraction efficiency.
- Exposed 14 REST routes across eight modules using **Python**, **SQLAlchemy**, and **PostgreSQL**, enabling a unified backend for Linq SMS and web clients.

HiivBuddy - 24x7 AI Advisory Board [\[link\]](#)

Jan 2026 - Feb 2026

- Developed an autonomous AI framework with five specialized roles and OpenAI-compatible LLM integration, designing the pipeline architecture and iterating implementation using Claude Code.
- Constructed the complete technical engine, reducing initial user-wait times to <200ms through non-blocking streaming data flow.
- Templatized a modular design, reducing feature addition time from several days to under 3 hours.

Mavenly - LLM Intelligence Pipeline [\[link\]](#)

Nov 2025 - Dec 2025

- Developed a backend service with LangChain and LangGraph that analyzes datasets from natural language and returns structured results with quality guardrails.
- Implemented **multi-agent ReAct workflows** that improved reasoning reliability across multiple risk categories. Created a Streamlit interface that allows non-technical users to get accurate insights quickly.

ZeroETL - Open Source High-Throughput Ingestion Library [\[link\]](#)

May 2025 - June 2025

- Designed and implemented a zero-ETL PySpark ingestion framework for Apache Iceberg with deduplication, snapshot management, and time-travel queries, enabling reliable and reproducible data states across batch pipelines.
- Open-sourced the library and improved ingestion reliability and operational efficiency for small-business data workflows by reducing manual processing overhead and improving pipeline performance and consistency with close to **820 downloads**.

Education

Arizona State University – MS in Computer Science(STEM)

Aug 2024 - May 2026

Cloud Computing, Statistical Machine Learning, Natural Language Processing

RV College Of Engineering– BE in Electrical Engineering

Aug 2020 - June 2024

Data Structures and Algorithms, OOP Principles, Machine Learning, RDBMS